

Assignment 2 — Build the loop

Agentic AI Workshop

Now build the agent, and find out whether it was worth it.

1. A LangGraph agent with at least four nodes — gather, assess, decide, act — and at least one conditional edge. Include the rendered graph.
2. A local model in the `assess` node. `gemma-3-1b-it`, or `gemma-3-4b-it` if the smaller one proves too unreliable. Its job is narrow: read log lines, classify the failure. It does not choose tools — your graph does. Say in your write-up why that split was necessary.
3. Two architectures, both built: a reflex version (rules only) and a deliberative or hybrid version (model in the loop). Same nodes where possible; the difference should be in the edges.
4. The comparison. Week 1's baseline, your reflex agent, and your model agent — on the same 100 alerts. Report accuracy, cost per alert, tool calls, and wall time.
5. Failure handling. Tools time out ~10% of the time. Show what your graph does about it.
6. Persistence. Add a checkpointer. Interrupt a run mid-alert and resume it. Show both.

Submit: notebook with the graph diagram, the four-way comparison table, and a paragraph on which architecture you would ship and why.

If your agent loses to the baseline, report that. A well-measured loss is worth more than an unmeasured win, and you will not be the only one.