

Assignment 3 — Make it shippable



Agentic AI Workshop

1. Calibration. Have your model emit a confidence with each classification. Measure whether that confidence means anything on 50 labelled alerts. Report ECE. Recalibrate if needed.
2. The escalation threshold. Derive p^* from your own cost numbers. Wire it into a conditional edge so low-confidence alerts route to a human.
3. The gate. `interrupt()` before every rollback. Your interrupt payload must show the evidence — which tools ran, what they returned, what could not be verified. Not a rationale.
4. Memory. Store outcomes across alerts. Has this service failed this way before? Show one case where memory changed the decision — and one where a wrong memory would have made things worse.
5. The harness. Iteration cap, token budget, wall-clock limit. Make at least two of them fire and log the stop reason.
6. Evaluation. Run on the held-out 100 alerts (`alerts[100:]`, which you have not tuned on). Report accuracy, cost per alert, and a 95% confidence interval on your accuracy.
7. Failure taxonomy. Categorise every failure from your traces. Counts, not adjectives.

Submit: the final agent, the eval results with the interval, and one honest paragraph: what would you need to measure before letting this run unattended on real alerts, and what do you still not know?

That paragraph carries the most credit in the whole project. Every demo works.